

Project 4: Bayesian Methods: Motion Estimation for Contrast Maximization with Event Cameras

AEROSP 567: Statistical Inference, Estimation, and Learning

Onur Bagoren

December 16, 2021

Abstract

Event cameras are biologically inspired cameras that, in contrast to traditional cameras, register changes in pixel intensity asynchronously. Each change in pixel intensity, referred to as an "event", is triggered by changes in motion in the scene at which the camera is observing; hence, events can be highly informative regarding the dynamics of the scene it is observing. The method highlighted in this project aims to solve the problem of motion parameter estimation of the ego-motion of the camera using probabilistic methods. Posing the problem of estimating motion parameters as an inference problem, this solution provides both information about the motion of the camera while also informing uncertainty of the estimations, which can later be used in tasks of computer vision in applications that require motion parameter estimates, such as scene segmentation or depth estimation. The proposed method utilizes available datasets online and aims to solve linear motion problems. This method proved to be successful in estimating motion parameters, yet suffered from computational complexity, making it unfavorable to use for online estimation problem, but favorable for post processing of the data.

1 Introduction

Event cameras are cameras that respond to pixel intensity changes in the pixel frame of the camera. An "event", occurs once the change of the pixel intensity is over 15%. Each event is represented in the spatio-temporal coordinates, along with a polarity value, which represents whether the pixel value increased, or decreased in intensity. Hence, each event registered by an event camera is represented as a tuple of these values $e_i = (t_i, x_i, y_i, p_i)$. In very simple terms, the product of an event camera can be analogous to

a point cloud, only with extra information about the direction at which an object is either approaching the camera or moving away from the camera. A visualization of a collection of events that are triggered by moving objects in front of an event camera is shown in fig. 1.




Figure 1: Events caused by moving objects in the scene of the camera

Event cameras additionally offer multiple advantages over traditional cameras. Event cameras have a high dynamic range, which measures up to 140 dB as compared to traditional cameras, which have a dynamic range of 60 dB. More importantly, event cameras offer very low latency, to the order of microseconds. A primary advantage of such low latency measurements prevent event camera observations to suffer from motion blurring. As the changes of pixel intensities also occur asynchronously, each individual object operating on an event camera can be registered at this low latency. These properties of event cameras can then be used for sample based motion estimation, rather than timing based motion estimation, as it is commonly done with traditional cameras that capture frames at certain time intervals. These advantages have lead to numerous advancements of the use

of event cameras in computer vision and image processing applications. As the event camera responds to the changes of objects in the scene it is observing, the dynamics of the scene can be captured. This is where the low latency especially helps, in that accurate representations of the objects that the camera is observing can be captured by the event camera.

A novel, and important, challenge regarding the use of event cameras have been tied to the concept of contrast maximization and parameter estimation with event cameras, which are two areas that are closely related for any event camera data to be used in any complex computer vision or image processing application. This project aims to utilize a probabilistic representation of the data and event generation of the event camera and utilize probabilistic methods in order to perform motion parameter estimation of the camera and observed objects. This challenge is referred to motion estimation. Contrast maximization is a resulting process after motion estimation, which acts as a filter to sharpen the image in order to reduce white noise produced by the sensor itself.

2 Motion Estimation

2.1 Representation of Motion with Event Cameras

Event cameras capture any change in pixel values, which can be seen as any motion that is captured in the frame at which the camera is observing. Hence, following the trace of the events over time is sufficient to characterize the motion parameters of the scene and camera.

Motion estimation with cameras are also analogous to optical flow estimation, which is the characterization of the rate of pixel movement across the image frame. This is typically the method at which such estimation tasks are completed, without information about the camera extrinsic parameters, which describe how the camera frame transforms into the world frame, it is difficult to estimate the actual rate of change.

The task for this project involves characterizing the motion parameters of linear movements of the camera. The model for characterizing such motion in the image coordinate frame is given with (2.1) and eq. (2.2).

$$\mathbf{x}(t_1) = \mathbf{x}(t_0) + \mathbf{v}_x t \quad (2.1)$$

$$\mathbf{y}(t_1) = \mathbf{y}(t_0) + \mathbf{v}_y t \quad (2.2)$$

The motion parameters are hence defined as $\Theta_x = [x(t_0) \ \mathbf{v}_x]^T$ and $\Theta_y = [y(t_0) \ \mathbf{v}_y]^T$.

2.1.1 Existing Solutions

The existing solutions [1], [2] for event camera motion estimation utilizes image processing techniques in order to both find the direction of highest event accumulation. Such solutions aim to group together, or warp, events so that a set of motion parameters Θ can describe the trajectory of the events. A warped event \mathbb{W} can then be characterized according to this definition shown in eq. (2.3):

$$\mathbf{x}'_k = \mathbb{W}(\mathbf{x}_k; t_k; \Theta) \quad (2.3)$$

$$\mathbf{x}_k = (x_k \ y_k)^T \quad (2.4)$$

where Θ represents a set of motion parameters in either the x or y direction of the pixel values. For optical flow estimation, the estimation parameters are computed for small periods of time, which allows representing the warped events as a linear equation, as shown in (2.5)

$$\mathbf{x}'_k = \mathbf{x}_k - (t_k - t_{ref})\theta \quad (2.5)$$

where t_{ref} is the first time at which the first event occurs. The warped events starting from t_{ref} are then projected to a discretized plane of pixels according to (2.6). This projection of events construct an image patch H , which encapsulates information regarding the set of events that fall within the trajectory constructed by the motion parameter θ . The general expression for the image patch is shown in eq (2.6):

$$H(\mathbf{x}; \theta) = \sum_{k=1}^{N_e} b_k \Delta(\mathbf{x} - \mathbf{x}'_k) \quad (2.6)$$

Where N_e is the number of events, $b_k = 1$ and $\Delta(\cdot)$ represents any function that evaluates the proximity of \mathbf{x}'_k to the actual events \mathbf{x}_k at time t_k . The selection of this $\Delta(\cdot)$ function is a design choice for the application in question.

Upon the design of this method of encapsulating the events in the image frame, an objective function to evaluate the selected parameters Θ is designed.

The objective function then becomes finding the parameter that maximizes the variance on the image patch H generated by θ , with eq. (2.7).

$$f_{\sigma^2}(\theta) = \frac{1}{N_p} \sum_{i,j} (h_{i,j} - \mu_H)^2 \quad (2.7)$$

where N_p is the number of pixels on H , $h_{i,j}$ represents the number of clustered events projected to pixel location (i, j) and μ_H is the mean number of clustered at a given pixel of H , such that

$$\mu_H = \frac{1}{N_p} \sum_{i,j} h_{i,j} \quad (2.8)$$

Equation (2.7) is one of the many objective functions that have been proven to be successful for these applications [2].

2.2 Proposed Solution to Motion Estimation

The proposed solution utilizes a probabilistic model to update the belief on the motion parameters as data from the event camera is received. This formulation poses the problem as a learning problem on the parameters.

As the dynamics model is linear, the problem unravels as a linear regression problem on the data in order to fit the parameters on the data. There are a multitude of algorithms that can be used to solve this problem. The solution with this project utilizes Gaussian Linear Models and poses it as a linear inverse problem to use Gaussian Processes to learn the parameters.

In order to formulate the problem, it is necessary to define the concepts that will be used in this problem formulation, and how it was used in the implementation of the proposed solution.

2.2.1 Distributions on the Variables

Unlike the solutions described in Section 2.1.1, where the motion parameters Θ_x and Θ_y were treated as deterministic parameters, the solution will approach the parameters as random variables, with probability distributions that correspond to the defined model.

The proposed solution The design choice for the probability distribution of the random variables were defined as normal distributions, where the distribution of parameters Θ_x and Θ_y are shown in eq. (2.9) and (2.10).

$$\mathbb{P}(\Theta_x) \sim \mathcal{N}(\mu_x, \Sigma_x) \quad (2.9)$$

$$\mathbb{P}(\Theta_y) \sim \mathcal{N}(\mu_y, \Sigma_y) \quad (2.10)$$

As part of the dynamic system, the manner in which the event camera captures information and registers events can also be posed as a probabilistic

model, due to the probabilistic distributions of the parameters provided in eq. (2.9) and eq.(2.10).

The measurement model for the event camera is a linear model [3], [4]. The generation of events are generated as an increase of contrast in the pixel relative to its previous contrast values, resulting in the intensity increase above a threshold that triggers an event. Hence, the the measurement model can be described as the linear model presented with eq. (2.11), where the measurement noise η is zero-mean with variance that depends on the event camera model that is used to take the measurements (more detail on event camera specifications are provided with [3], Table 1).

$$Y = A\Theta + \eta \quad (2.11)$$

$$\eta \sim \mathcal{N}(0, \Gamma) \quad (2.12)$$

$$A = \begin{bmatrix} 1 & t^{(1)} \\ \vdots & \\ 1 & t^{(N)} \end{bmatrix} \quad (2.13)$$

In eq. (2.11), the matrix A represents the Vandermonde matrix, which has dimensions of $\mathbb{R}^{N \times 2}$, where N is the number of data points that are measured at a time t and 2 is for the number of parameters of Θ . As there are motion parameters to learn in both the x and y direction, the measurement model is then regarded as separate in both directions.

Provided this information, the conditional probabilities between the two random variables representing the measurements and the motion parameters can then be constructed. The measurement model implies a linear Gaussian relation between the parameters due to the Gaussian distribution on the parameters in (2.9) and (2.10); hence, the distribution shown in eq. (2.14) can be constructed.

$$\mathbb{P}(Y|\Theta) \sim \mathcal{N}(\Theta t, \Gamma) \quad (2.14)$$

2.2.2 Linear Gaussian Models

The proposed solution to the problem of motion parameter estimation utilizes a linear Gaussian model. This proposal stems from the idea that the prior distribution on the parameters are Gaussian, as shown in Section 2.2.1, along with the likelihood model being a linear measurement model. From conjugate pairs, it is then known that with a linear likelihood model, the posterior will be of the same distribution family as the prior; in this case a Gaussian distribution, which belongs to the exponential distribution family.

The linear Gaussian model utilizes Bayesian Inference in order to learn the posterior distribution of

a parameter, provided an initial distribution of the prior and a given likelihood model.

Bayesian Inference For Bayesian Inference, the learning is done by updating the belief on the parameters θ after observing the data. This belief is a continuous representation that can make predictions between the data that was seen. In order to update this belief representation on the parameters after observing a set of data, Bayes' rule can be used.

$$\mathbb{P}(\Theta|Y) = \frac{\mathbb{P}(Y|\Theta)\mathbb{P}(\Theta)}{\mathbb{P}(Y)} \propto \mathbb{P}(Y|\Theta)\mathbb{P}(\Theta) \quad (2.15)$$

In eq. (2.15),

1. $f_{\Theta|D}(\theta|d)$ is the posterior distribution which represents the distribution of the belief on parameters θ after seeing data d . In the scope of this project, the prior distributions are represented with equations (2.9) and (2.10).
2. $f_{D|\Theta}(d|\theta)$ is the sampling distribution, which represents how the data will be interpreted by the learning model. The sampling distribution is key in how the belief is represented continuously after observing data, as it varies from model to model. Will be represented as $\mathcal{L}(\theta)$. The likelihood model in this project is constructed from the measurement model, which is shown with eq. (2.11)
3. $f_{\Theta}(\theta)$ is the prior distribution, sometimes shown as $f_{\Theta}(\theta|I)$ where I represents the available information. The effectiveness of Bayesian inference depends on the interpretation of I to construct the prior distribution. A good analogy for the importance of the prior is that learning with existing beliefs about the event/topic will impact how the learning will occur with new information
4. $f_D(d)$ is the evidence, which will commonly be used as a normalization factor η as it can be very difficult to compute.

Useful properties of Bayesian Inference:

1. Completing inference on data sequentially or batch does not impact the inference

Proof.

$$p(\theta|A, B) = p(\theta|A) \frac{p(B|A, \theta)}{p(B|A)} \quad (2.16)$$

$$= p(\theta) \frac{p(A|\theta)}{p(A)} \frac{p(B|A, \theta)}{p(B|A)} \quad (2.17)$$

$$= p(\theta) \frac{p(A, B|\theta)}{p(A, B)} \quad (2.18)$$

$$= p(\theta|A, B) \quad (2.19)$$

□

The linear Gaussian problem utilizes the fact that the joint distribution between the parameters and the measurement are jointly Gaussian. To show this, it must first be established what the marginal distribution of the random variable for the measurement will be. As the measurement model from eq. (2.11) is just scaling and addition of Gaussian RVs, the marginal distribution for the measurement becomes the expression shown with (2.20).

$$\mathbb{P}(Y) \sim \mathcal{N}(A\Theta, A\Sigma A^T + \Gamma) \quad (2.20)$$

$$\text{Cov}[Y, \Theta] = A\Sigma \quad (2.21)$$

With the marginal distribution of both Θ and Y established, their joint distributions then are formed as (2.22), as their covariance can be expressed as shown in (2.21). The joint distribution is then normally distributed.

$$\begin{bmatrix} \Theta \\ Y \end{bmatrix} \sim \mathcal{N} \left(\begin{bmatrix} \mu_{\Theta} \\ A\mu_{\Theta} \end{bmatrix}, \begin{bmatrix} \Sigma_{\Theta} & \Sigma_{\Theta} A^T \\ A\Sigma_{\Theta} & A\Sigma_{\Theta} A^T + \Gamma \end{bmatrix} \right) \quad (2.22)$$

Finally, in order to get the posterior distribution $\mathbb{P}(\Theta|Y)$, it can be seen that conditioning the parameters on the data will be sufficient. For jointly distributed Gaussian RVs, the conditioning of the variables upon each other yield a closed form expression shown with equations (2.24) and (2.25)

$$\mathbb{P}(\Theta|Y) \sim \mathcal{N}(\mu_{\Theta}^{post}, \Sigma_{\Theta}^{post}) \quad (2.23)$$

$$\mu_{\Theta}^{post} = \mu_{\Theta} + \Sigma_{\Theta} A^T (A\Sigma_{\Theta} A^T + \Gamma)^{-1} (y - A\mu_{\Theta}) \quad (2.24)$$

$$\Sigma_{\Theta}^{post} = \Sigma_{\Theta} - \Sigma_{\Theta} A^T (A\Sigma_{\Theta} A^T + \Gamma)^{-1} A\Sigma_{\Theta} \quad (2.25)$$

where y is the data.

Utilizing the closed form expression of the Linear Gaussian, these equations can then be used to perform linear regression on the motion parameters.

2.3 Implementation of the Linear Gaussian Inverse Problem

2.3.1 Partitioning the data

Following the assumptions that are made regarding optical flow estimation, at a short enough time frame, all representations of the object motion can be linearized. Due to the advantage of the extremely high sampling rate of the event camera, within small bursts of times, there is a large amount of data to perform inference on.

Taking advantage of the structure of the data, the data is then segmented into small portions, each portion representing an area to run the regression on to learn the parameters of that time slice. This offers two solutions:

1. Segmenting and operating on less data implies that the size of the Vandermonde matrix, which is inverted with in the GP equations for the posterior moment ($\mu_{\Theta}^{post}, \Sigma_{\Theta}^{post}$) computations in eq. (2.24) and eq. (2.25), is smaller. As the inversion of a matrix runs at the time complexity of $\mathcal{O}(n^3)$, the high data regime runs the risk of the algorithm not running fast enough to compute motion parameters within a reasonable time.
2. Transferring this solution to an online problem can then yield the advantage performing the computation at timestamps.

In addition to mitigating the size of the matrices to be inverted, there are also other ways to mitigate with inverting a large dimensional matrix.

Woodbury Matrix Identity is an identity that follows the proposition that

$$(A + UCV)^{-1} = A^{-1} + A^{-1}U(C^{-1} + VA^{-1}U)^{-1}VA^{-1}$$

This identity can be used in equations (2.24) and (2.25). These equations then become

$$\mathbb{P}(\Theta|Y) \sim \mathcal{N}(\mu_{\Theta}^{post}, \Sigma_{\Theta}^{post}) \quad (2.26)$$

$$\mu_{\Theta}^{post} = (A^T \Gamma^{-1} A + \Sigma_{\Theta}^{-1})^{-1} (A^T \Gamma^{-1} y + \Sigma_{\Theta}^{-1} \mu_{\Theta}) \quad (2.27)$$

$$\Sigma_{\Theta}^{post} = (A^T \Gamma^{-1} A + \Sigma_{\Theta}^{-1})^{-1} \quad (2.28)$$

The advantage of this representation is that rather than inverting a matrix of the dimension of the data, a matrix of the dimension of the parameters is being inverted - vastly saving computation data for the case

of event camera applications, which is a dense data application with few parameters.

For the implementation in the project, there partitioning was done every 7500 data points. For the dataset that was used, this yielded 100 segments to run inference on.

2.3.2 Understanding the form of the data

A large challenge about working with event cameras come in multiple forms. The primary issue is in the form of asynchronous data generation. Asynchronous data generation causes the data at different time steps to be different sizes. In comparison to more traditional sensors, which query for a signals at defined time intervals, the event camera will output an event whenever there is any motion. This poses the problem with working with multidimensional data for the task of estimating a single parameters. Figure 2 displays this phenomenon, that over a time period, multiple points can be present.




Figure 2: Data generated within a small time window. this figure shows that at different time steps, there are multiple dimensions of data, all representing the same information about the position of an object on the x- or y-axis.

This issue becomes prevalent when forming the Vandermonde matrix in the measurement model setup. The Vandermonde matrix is a representation for the linear transformation of the motion parameters over time. The nature of the a synchronicity

of the data then requires that the matrix be varied-dimensional across different time values. The necessity to then iterate over each time step to obtain yields an iterative/sequential solution in comparison to the batch solution.

Shown with eq. (2.19), for Bayesian inference, sequential vs. batch does not make a difference for the inference that is being done, but in terms of runtime, iterative algorithms in comparison to the batch algorithms suffer from the not being vectorized, which greatly improves the computation speed of processes.

Two solutions were attempted in order to mitigate this problem.

1. **Representing ragged data as a matrix with fixed size:** This solution required pre-processing the data to represent it as a matrix with fixed size. The matrix would be of dimension $d \times n$, where d represents the number of unique timestamps that the event camera measured events on, and n representing the maximum number of events that were triggered at a single time step. The entries where the number of data-points $< n$ are then represented as "not a number".
2. **Averaging the measurements:** This method mitigates the varied-dimensional nature of the data by rather than evaluating the likelihood at each data point, evaluating at the mean of the measurements. The assumption made for this solution follows the assumption that due to the linear measurement model and the operation of calculating being a linear operation, the evaluation of the likelihood at a certain time would be equivalent to the measurement of the mean of the measurements at that same time.

For the implementation of the project, the second method was used, as it required less offline time in order to pre-process the data and form a new representation of the data, along with an ease of implementation, as single dimensional data can be input into the measurement model without adjusting dimensional over time and event segment (discussed in Section 2.3.1)

2.4 Converting from camera frame to image frame

The data that is being processed with this algorithm outputs event positions on the image frame. The image frame can be represented as $\mathbf{u} = [u \ v]^T$, where u is the pixels in the x-direction and v is the pixels

in the y-direction. The result of running the computations on the pixel frame is the lack of ability to translate the rate of change to more intuitive measures, such as m/s .

A solution for this comes in the form of identifying and knowing the characteristics of the camera that is being used. These are known as the intrinsic parameters of the camera. The intrinsic parameters convey information regarding the focal length in the x- and y-direction, along with the principle point of the camera coordinates (principal point is the offset of the center of the camera frame to the top left corner of the frame, measured in).

By representing the pixel coordinates as homogeneous coordinates, the projection of the pixel coordinates can then be done on the world coordinates. This method follows eq. (2.29), where $\tilde{\mathbf{x}}_c$ are the world coordinates in 3D and $\tilde{\mathbf{u}}$ are the homogeneous image coordinates.

$$\tilde{\mathbf{u}} = M_{int}\tilde{\mathbf{x}}_c \quad (2.29)$$

For this project, the intrinsic values of the camera were provided with the test dataset, which were used to compute the motion parameters to the image plane to compare results with the predictions.

2.5 Informative Priors

As discussed in Section 2.2.2, a good prior is crucial for the inference problem. In the implementation of this project, the partitioning of the data was utilized in order to design informative priors in between segments.

As it is assumed that each partition is the continuation of the partition before it, the prior distribution on the parameters for the GP on the next segment was selected according to the prior before it.

The prior set for the initial segment is $\mathbb{P}(\Theta) \sim \mathcal{N}(0, I)$ for both the x-direction and the y-direction. Selecting this distribution as the prior for each partition will cause issues with the predictions, as assuming not only zero velocity, but also a zero-position at each time step will cause any change in position to impact both parameters at a large scale, yielding inaccurate predictions.

2.6 Confidence Interval Computation

A method to quantify the confidence in the estimation is through a 95% Confidence Interval (CI). In order to evaluate the 95% CI, 10E5 samples from the posterior distribution at segments were taken. Then,

a Monte Carlo estimator was used in order to estimate the motion parameters. The setup is as shown in eq. (2.30). In this equation 1_x is the indicator function, which is 1 when the sample generated from the MC estimator is equal to that of the mean of the posterior, otherwise it is 0.

$$S_n = \mathbb{E}[g(x)] = \frac{1}{n} \sum_{i=0}^{10E5} 1_x \quad (2.30)$$

$$1_x = \begin{cases} 1 & \text{if } \mu_{pred} = \mu_{estimate} \\ 0 & \text{otherwise} \end{cases} \quad (2.31)$$

The CI bounds are then computed as shown in eq. (2.32), where $\frac{\sigma}{\sqrt{n}}$ is the standard error of the MC estimator.

$$\mathbb{P}\left(S_n - \frac{z\sigma}{\sqrt{n}} \leq \mu_\Theta \leq S_n + \frac{z\sigma}{\sqrt{n}}\right) = 0.95 \quad (2.32)$$

3 Results

For the testing of the algorithm, the data representing linear motion from the event camera dataset was used [5]. The dataset was produced from measurements taken with a DAVIS240C event camera, which from Table 1 in [3] can be seen as having measurement noise of 0.1 events per pixel per second.

The results section will analyze the inference run on 4 of the segments. These segments are the first segment, 50th segment and last segment. These were selected in order to display the evolution of the predictions over time, and the impact of informative priors.

For visual familiarity with the environment that the event camera is interacting with, a grayscale image of the environment that the event camera is observing is provided with fig. 3.




Figure 3: Grayscale image of the environment the event camera observes

3.1 Evaluating the Results

The benchmarking of the performance of the algorithm was done against ground truth data, which is provided with the data upon which the algorithm was tested with. The ground truth data for the dataset "slider_{close}" was computed to be $(46.96, 0) \frac{px}{s}$ for the camera motion. This computation was done using eq. (2.29).

3.1.1 First Partition

This section provides figures regarding the posterior distribution after the inference, the prediction in comparison to the ground truth and a visual representation of what the prediction implies for working with the data.

Visualization of the prediction The visualization for interpreting what this prediction implies is helpful to get context on the problem that is trying to be solved. The collection of events are shown in 4. The red line in fig 4 represents the prediction made regarding the trajectory of the events over time, while also being a concise representation of the motion parameters.




Figure 4: Events generated in the first segment

Figure 5 displays a figure for which the frame aligns with the predicted motion.

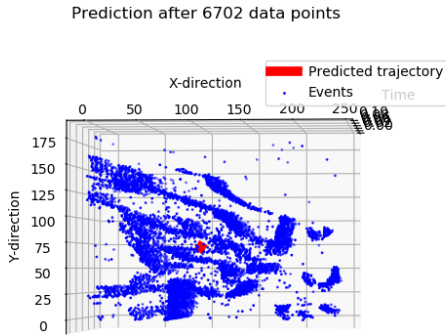


Figure 5: Events aligned according to the prediction made by the algorithm, shown in a 3D representation

A more informative can be seen with figure 6, where the predicted trajectory are shown in 2D plots, for each direction in x and y.

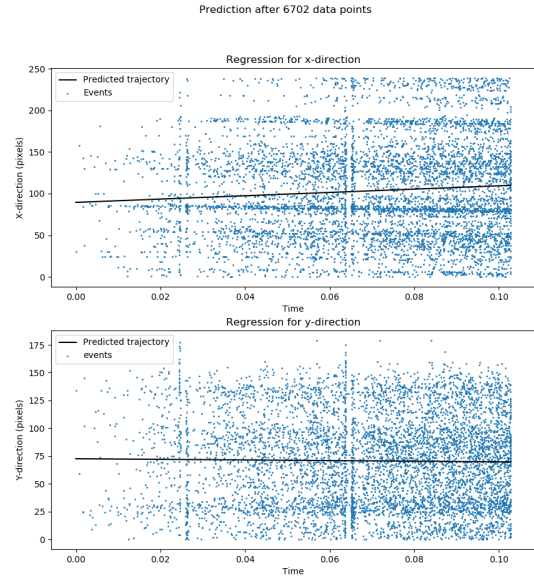


Figure 6: Events aligned according to the prediction made by the algorithm, shown in each dimension of x and y

Ground Truth Comparison The comparison of the predicted velocity from the motion parameters, is compared to that of the ground truth data. As the ground truth data does not provide information regarding the initial position, only the comparison between the ground truth velocity can be done.

Figure 7 displays the predicted velocity against the true velocity at the first segment.

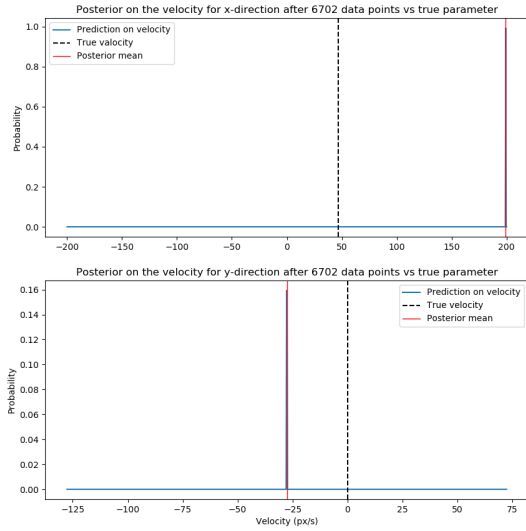


Figure 7: The comparison of the predicted trajectory to the ground truth

Confidence Interval For the first segment, the 95% confidence bounds for the velocity estimate are shown in eq. (3.2)

$$\text{CI for x-direction: } 199.205 \leq 199.205 \leq 199.205 \quad (3.1)$$

$$\text{CI for y-direction: } -27.626 \leq -27.626 \leq 27.626 \quad (3.2)$$

3.1.2 50th Partition

Visualization of the prediction The visualization for interpreting what this prediction implies is helpful to get context on the problem that is trying to be solved. The collection of events are shown in 8. The red line in fig 8 represents the prediction made regarding the trajectory of the events over time, while also being a concise representation of the motion parameters.

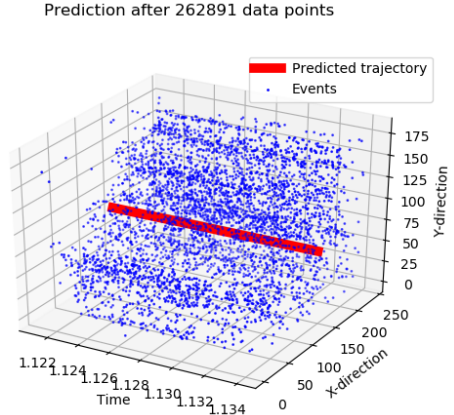


Figure 8: Events generated in the first segment

Figure 9 displays a figure for which the frame aligns with the predicted motion.

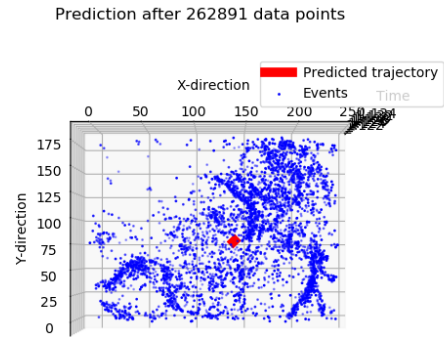


Figure 9: Events aligned according to the prediction made by the algorithm, shown in a 3D representation

A more informative can be seen with figure 10, where the predicted trajectory are shown in 2D plots, for each direction in x and y.

Ground Truth Comparison The comparison of the predicted velocity from the motion parameters, is compared to that of the ground truth data. As the ground truth data does not provide information regarding the initial position, only the comparison between the ground truth velocity can be done.

Figure 11 displays the predicted velocity against the true velocity at the first segment.

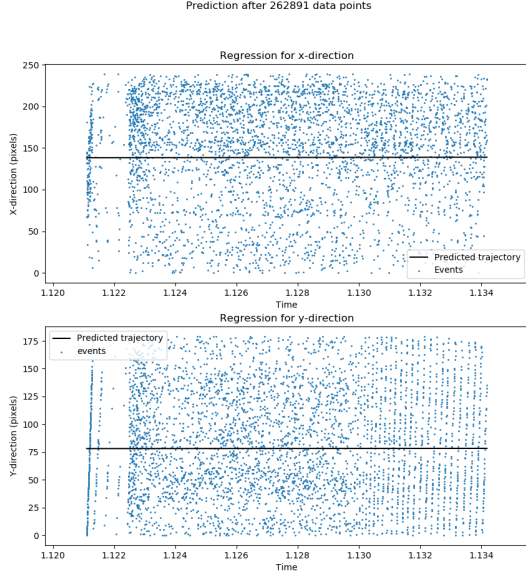


Figure 10: Events aligned according to the prediction made by the algorithm, shown in each dimension of x and y

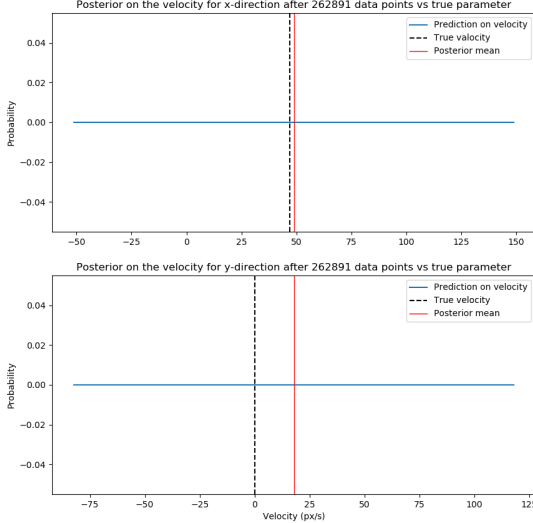


Figure 11: The comparison of the predicted trajectory to the ground truth

Confidence Interval A method to quantify the confidence in the predictions is to use confidence intervals. For the 50th segment, the 95% confidence bounds for the velocity estimate are shown in eq.

$$(3.4)$$

$$\text{CI for x-direction: } 48.816 \leq 48.816 \leq 48.816 \quad (3.3)$$

$$\text{CI for y-direction: } 17.802 \leq 17.801 \leq 17.801 \quad (3.4)$$

3.1.3 100th Partition

Visualization of the prediction The visualization for interpreting what this prediction implies is helpful to get context on the problem that is trying to be solved. The collection of events are shown in 12. The red line in fig 12 represents the prediction made regarding the trajectory of the events over time, while also being a concise representation of the motion parameters.

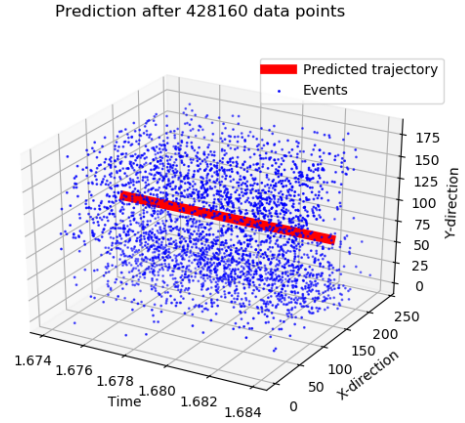


Figure 12: Events generated in the first segment

Figure 13 displays a figure for which the frame aligns with the predicted motion.




Figure 13: Events aligned according to the prediction made by the algorithm, shown in a 3D representation

A more informative can be seen with figure 14, where the predicted trajectory are shown in 2D plots, for each direction in x and y.

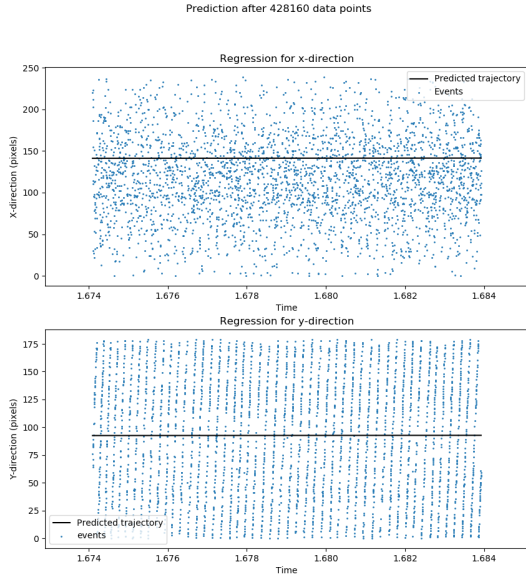


Figure 14: Events aligned according to the prediction made by the algorithm, shown in each dimension of x and y

Ground Truth Comparison The comparison of the predicted velocity from the motion parameters, is compared to that of the ground truth data. As the ground truth data does not provide information regarding the initial position, only the comparison between the ground truth velocity can be done.

Figure 15 displays the predicted velocity against the true velocity at the first segment.

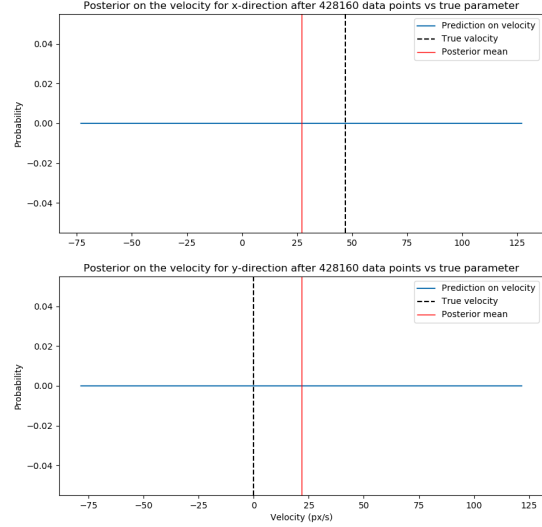


Figure 15: The comparison of the predicted trajectory to the ground truth

Confidence Interval A method to quantify the confidence in the predictions is to use confidence intervals. For the 100th segment, the 95% confidence bounds for the velocity estimate are shown in eq. (3.6)

$$\text{CI for x-direction: } 26.929 \leq 26.929 \leq 26.929 \quad (3.5)$$

$$\text{CI for y-direction: } 21.619 \leq 21.619 \leq 21.619 \quad (3.6)$$

3.1.4 Overall Performance

In order to evaluate the overall performance of the parameters that were learned, the posteriors of the parameters at each partition were plotted over time to observe convergence characteristics, shown with fig. 16.




Figure 16: Estimations of velocity over each segment plotted against the true velocity in pixels/s

4 Discussion

4.1 Computation Evaluation

The most drastic performance issue that arose from this method came from the computational complexity. Despite using the Woodbury Matrix Identity to reduce computational complexity with inverting large matrices, or pre-processing matrices with known structures to leverage computational processing, inverting Σ_{Θ} , which was a $\mathbb{R}^{7500 \times 7500}$ matrix slowed the computation of each segment to 3 seconds each. This would cause drastic issues in online computations, as it would not be capable of processing data at a high enough speed to be viable in practice.

4.2 Evaluation of Inferred Parameters

4.2.1 Performance on Partitioned Data

The first partition displays the necessity to have informed priors. It can be seen from both the visuals with fig. 6 that the trajectory does not agree with the trend of the data, along with the comparison to the ground truth value, shown with 7.

Although this impacted the prediction for the first partition, this was quickly corrected for in preceding partitions. Using informed priors that utilize the posterior distribution from the previous inference step prevented such issues from occurring for future partitions.

This is especially evident in the 50th and 100th partition estimates, where the estimates were closer to

the actual trend of the data and also closer to the ground truth data.

Reduction on the covariance of the estimations

A noticeable effect throughout the inference steps was the rapid drop in the uncertainty associated with the estimations. This was apparent in the 95% CI of the estimations, there the lower bound and upper bound were the same value up to the 10⁻⁴th decimal. A driving reason for this occurs from the fact that the abundance of data for the computation. The only method in which the uncertainty for the estimates to drop is in eq. (2.25).

The implication of the low uncertainty, yet inaccurate prediction at certain intervals point to the necessity to re-evaluate the inference model. Although over time the predictions agreed with the ground truth data, improvements to provide more interpretable and reliable predictions must be developed.

4.2.2 Performance on the entire dataset

Figure 16 represents the convergence of the motion parameters across the estimates. It can be seen that especially for the velocity estimate in the x-direction, which is the direction where there was the linear motion, converges after 40 partitions. It is also relevant to note that the steady state error associated state estimations on the parameters account for ≈ 20 pixels/s.

5 Conclusion

The learning method that utilized the linear-Gaussian method proved to be effective for offline computation. As the event camera provides an abundant amount of data over short periods of time, a probabilistic method in order to estimate the motion parameters is successful in reducing the uncertainty with the estimation, along with providing an understandable interpretation of the data.

The downside of the method is the computational complexity and the interaction with the structure of the data. There must be assumptions made about the linearity of the measurement model in order to implement the method described in this project. If this assumption is not made, the idea to take the average of the measurements to compute the likelihood cannot be done, and a more computational expensive and difficult to implement method must be taken.

The main task highlighted by the work for this project was to understand the data that is being worked with and model the solution to work best with the data. This became prevalent when working with the abundant and highly unstructured data that the event camera output. Developing methods to segment the data and take advantage of the underlying structure of the data not only lead to interpretable outputs regarding the predictions; but also, provide computational ease for the method as well.

References

- [1] G. Gallego, H. Rebecq, and D. Scaramuzza, “A unifying contrast maximization framework for event cameras, with applications to motion, depth, and optical flow estimation,” *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2018.
- [2] T. Stoffregen and L. Kleeman, “Event Cameras, Contrast Maximization and Reward Functions: An Analysis,” in *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. Long Beach, CA, USA: IEEE, Jun. 2019, pp. 12 292–12 300. [Online]. Available: <https://ieeexplore.ieee.org/document/8954356/>
- [3] G. Gallego, T. Delbruck, G. M. Orchard, C. Bartolozzi, B. Taba, A. Censi, S. Leutenegger, A. Davison, J. Conradt, K. Daniilidis, and D. Scaramuzza, “Event-based Vision: A Survey,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, pp. 1–1, 2020. [Online]. Available: <https://ieeexplore.ieee.org/document/9138762/>
- [4] P. Lichtsteiner, C. Posch, and T. Delbruck, “A 128×128 120 db 15 μ s latency asynchronous temporal contrast vision sensor,” *IEEE Journal of Solid-State Circuits*, vol. 43, no. 2, p. 566–576, 2008.
- [5] E. Mueggler, H. Rebecq, G. Gallego, T. Delbruck, and D. Scaramuzza, “The event-camera dataset and simulator: Event-based data for pose estimation, visual odometry, and slam,” *The International Journal of Robotics Research*, vol. 36, no. 2, p. 142–149, 2017.
- [6] —, “The event-camera dataset and simulator: Event-based data for pose estimation, visual odometry, and slam,” *The International Journal of Robotics Research*, vol. 36, no. 2, p. 142–149, 2017.

1 Appendix

1.1 Prior and Posterior Distributions at Partitions

1.1.1 First Partition

Prior Distribution for the first segment was selected as the standard Gaussian $\mathcal{N}(, I_{2 \times 2})$.

The prior on the motion parameters are shown with figures 1 and 2.

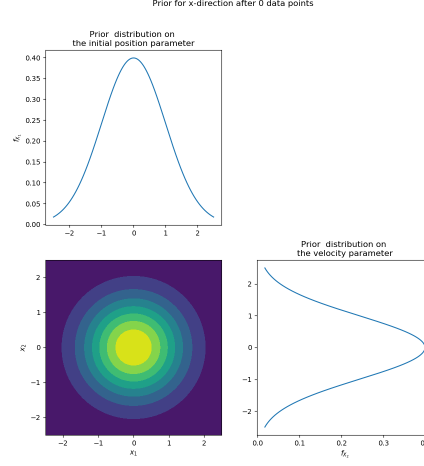


Figure 1: Prior distribution on the motion parameters in the x-direction at the first segment of the data

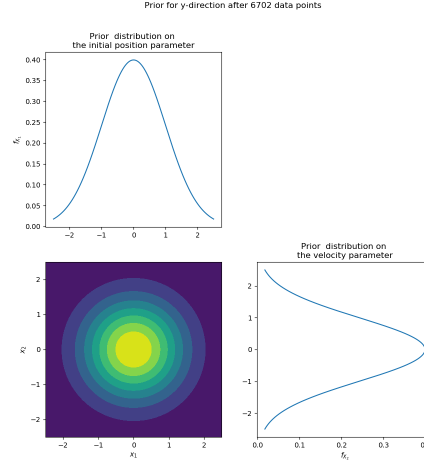


Figure 2: Prior distribution on the motion parameters in the y-direction at the first segment of the data

Posterior Distribution for the first segment is shown in figures 3 and 4.

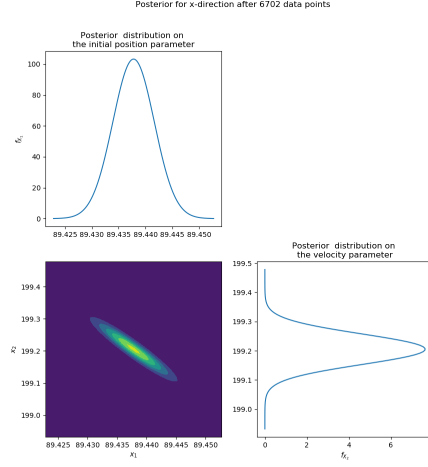


Figure 3: Prior distribution on the motion parameters in the x-direction at the first segment of the data

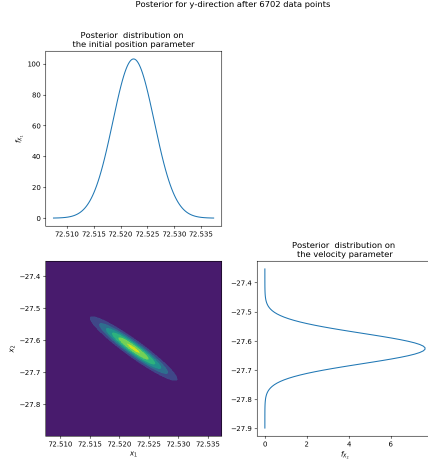


Figure 4: Prior distribution on the motion parameters in the y-direction at the first segment of the data

1.1.2 50th Partition

Prior Distribution for the first segment was selected as the standard Gaussian $\mathcal{N}(, I_{2 \times 2})$. The prior on the motion parameters are shown with figures 5 and 6.

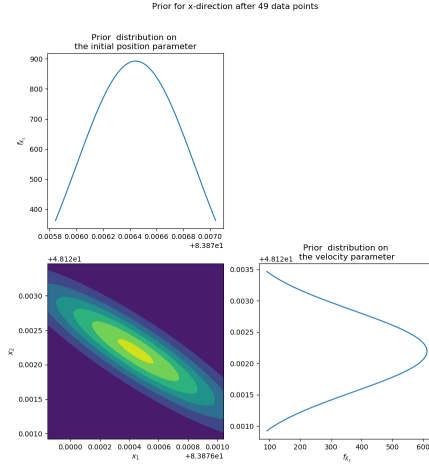


Figure 5: Prior distribution on the motion parameters in the x-direction at the 50th segment of the data

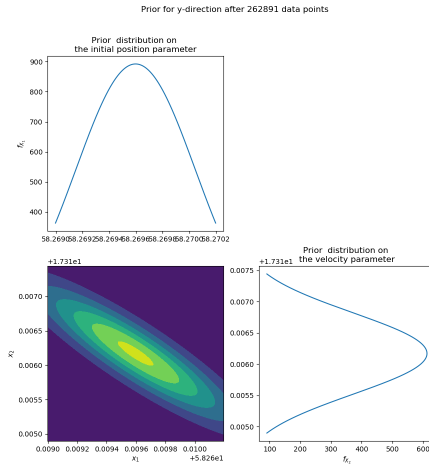


Figure 6: Prior distribution on the motion parameters in the y-direction at the 50th segment of the data

Posterior Distribution for the first segment is shown in figures ?? and ??.

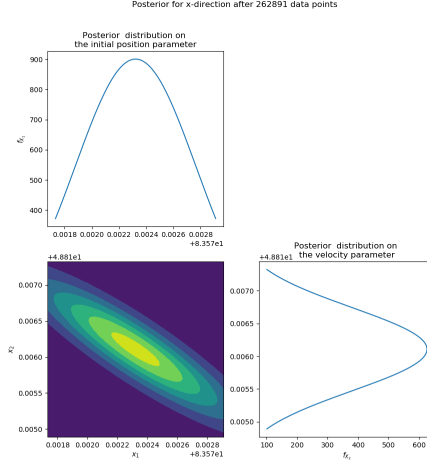


Figure 7: Prior distribution on the motion parameters in the x-direction at the 50th segment of the data

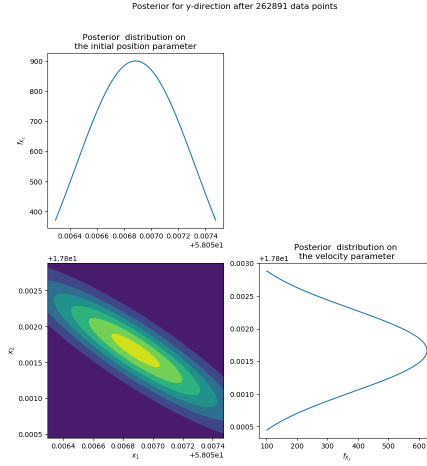


Figure 8: Prior distribution on the motion parameters in the y-direction at the 50th segment of the data

1.1.3 100th Partition

Prior Distribution for the first segment was selected as the standard Gaussian $\mathcal{N}(\cdot, I_{2 \times 2})$.

The prior on the motion parameters are shown with figures 9 and 10.

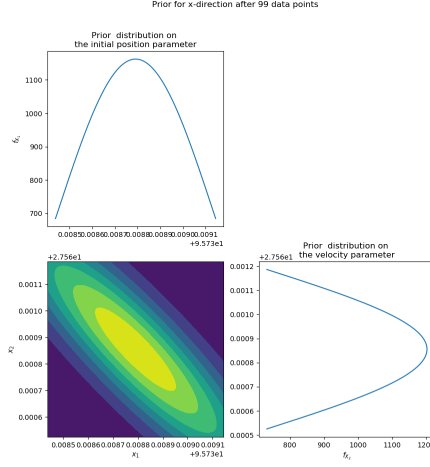


Figure 9: Prior distribution on the motion parameters in the x-direction at the 100th segment of the data

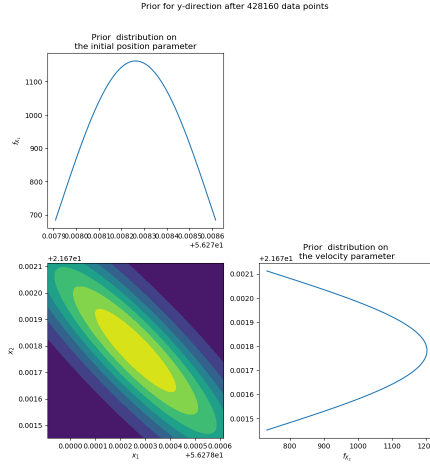


Figure 10: Prior distribution on the motion parameters in the y-direction at the 100th segment of the data

Posterior Distribution for the first segment is shown in figures 11 and 12.

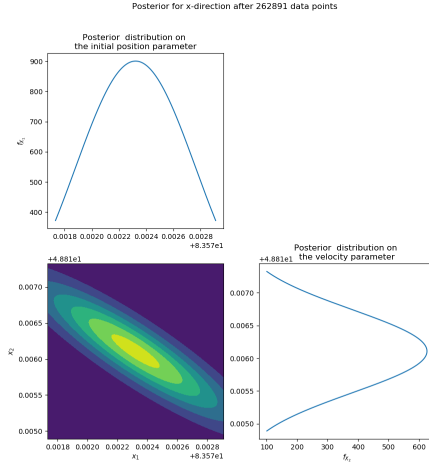


Figure 11: Prior distribution on the motion parameters in the x-direction at the 100th segment of the data

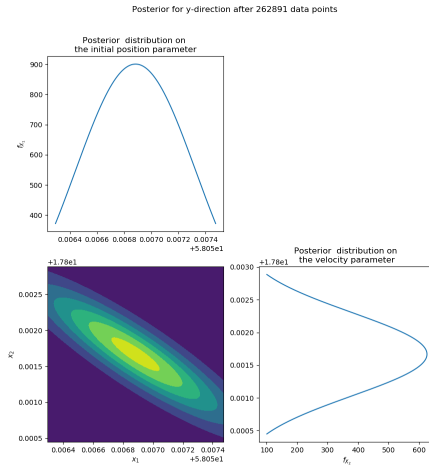


Figure 12: Prior distribution on the motion parameters in the y-direction at the 100th segment of the data

1.2 Code

1.2.1 Linear Gaussian Regression

```

1 import glob
2 from operator import mul
3 import os
4 import sys
5
6 import matplotlib.pyplot as plt
7 from mpl_toolkits.mplot3d import Axes3D
8 import numpy as np
9 import scipy.stats as stats
10 import target_distributions as td
11 from event_plotter import EventPlotter
12 from target_distributions import evaluate_gaussian
13 from plotter import plot_bivariate_gauss
14

```

```

15 from tqdm import tqdm
16
17
18 def compute_gaussian_predicted_marginal(mu, Sigma, A, Gamma):
19     """Compute the marginal distribution of Y where
20
21      $Y = A \theta + \xi$ ,  $\xi \sim N(0, \text{Gamma})$ 
22
23     with  $\theta$  distributed according to a Gaussian with mean  $\mu$  and
24     covariance  $\text{Sigma}$ 
25
26      $\theta$  is a vector of size (d)
27
28     Inputs
29     -----
30     mu : (d, ) array, the mean of  $\theta$ 
31     Sigma: (d, d) array, the covariance of  $\theta$ 
32     A : (n, d) linear model
33     Gamma: (n, n) array, the covariance of  $\xi$ 
34
35     Returns
36     -----
37     (mean, covariance, ASigma)
38     """
39     ASigma = np.dot(A, Sigma)
40     cov = np.dot(ASigma, A.T) + Gamma
41     return np.dot(A, mu), cov, ASigma
42
43
44 def linear_gaussian_inverse_problem(
45     prior_mean, prior_cov, A, Gamma, data, inverted_noise_var=False
46 ):
47     """Compute the posterior of a linear-Gaussian inverse problem
48
49     Inputs
50     -----
51     prior_mean: (d, ) prior mean
52     prior_cov : (d, d) prior Covariance
53     A : (M, d) linear model
54     Gamma : (M, M) noise model
55     data : (M, ) data
56
57     Outputs
58     -----
59     (post_mean, post_covcov): posterior mean and posterior covariance
60     """
61     # Check if all entries for gamma are equal
62     inverted = False
63     if np.all(np.diag(Gamma) == Gamma[0, 0]):
64         inverted = True
65         Gamma = np.eye(Gamma.shape[0]) * 1 / Gamma[0, 0]
66
67     # Compute Marginal
68     mean_pred, cov_pred, ASigma = compute_gaussian_predicted_marginal(
69         prior_mean, prior_cov, A, Gamma
70     )
71
72     if inverted:
73         print(np.linalg.inv(cov_pred))
74         first_expression = A.T @ Gamma @ A + np.linalg.inv(cov_pred)
75         second_expression = A.T @ Gamma @ data + np.linalg.inv(cov_pred) @ prior_mean
76         post_mean = np.linalg.solve(first_expression, second_expression)
77         post_cov = np.linalg.inv(first_expression)
78     else:
79         print(A.T @ np.linalg.solve(Gamma, A))
80         post_mean = np.linalg.solve(

```

```

81         A.T @ np.linalg.solve(Gamma, A) + np.linalg.inv(prior_cov),
82         A.T @ np.linalg.solve(Gamma, data) + np.linalg.solve(prior_cov, prior_mean),
83     )
84     post_cov = np.linalg.inv(
85         A.T @ np.linalg.solve(Gamma, A) + np.linalg.inv(prior_cov)
86     )
87     return post_mean, post_cov
88
89
90 def linear_regression(
91     xtrain, ytrain, xpredict, prior_mean, prior_cov, noise_var, inverted_noise_var=False
92 ):
93     """Perform linear regression
94
95     Inputs
96     -----
97     xtrain      : (N, d) training inputs
98     ytrain      : (N) training outputs
99     xpredict    : (M, d) an array of points where predictions are required (If none, then is
100                   ignored)
101     prior_mean  : (d+1) prior mean
102     prior_cov   : (d+1, d+1) prior covariance
103     noise_var   : (N, ) noise variance of each individual output
104
105     Outputs
106     -----
107     post_mean   : posterior mean of the parameters
108     post_cov    : posterior covariance of the parameters
109     pred_mean   : mean prediction at the predict points (optional)
110     pred_cov    : covariance of the prediction (optional)
111
112     Notes
113     -----
114     Predictions assumes no noise: just  $y = x^T \theta$ 
115     x
116     """
117     # Perform inference
118     N, d = xtrain.shape
119     A = np.ones((N, d + 1))
120     A[:, 1:] = xtrain
121     # Gamma = np.diag(noise_var)
122     post_mean, post_cov = linear_gaussian_inverse_problem(
123         prior_mean,
124         prior_cov,
125         A,
126         noise_var,
127         ytrain,
128         inverted_noise_var=inverted_noise_var,
129     )
130
131     # Perform prediction: which is just computing the marginal with a different
132
133     if xpredict is not None:
134         Npred, dpred = xpredict.shape
135         assert d == dpred, "prediction and testing features must be the same"
136         Apred = np.ones((Npred, d + 1))
137         Apred[:, 1:] = xpredict
138         pred_mean, pred_cov, _ = compute_gaussian_predicted_marginal(
139             post_mean, post_cov, Apred, 0
140         )
141         return post_mean, post_cov, pred_mean, pred_cov
142     else:
143         return post_mean, post_cov
144
145 def store_events_in_matrix(filename):

```

```

146     """
147     Store the events in a matrix.
148
149     Inputs
150     =====
151     events: numpy array
152     num_of_repeat: numpy array
153     """
154     EVENT_FILE = (
155         f"{sys.path[0]}/../data/new_slider_close/slider_close/split_data/{filename}"
156     )
157     events = np.genfromtxt(EVENT_FILE, delimiter=" ")
158     t = events[:, 0]
159     x = events[:, 1]
160     y = events[:, 2]
161     p = events[:, 3]
162
163     uniquew, inverse = np.unique(t, return_inverse=True)
164     counts = np.bincount(inverse)
165     max_count = np.max(counts)
166
167     uniques_x = np.zeros(shape=(uniquew.shape[0], max_count + 3))
168     uniques_y = np.zeros(shape=(uniquew.shape[0], max_count + 3))
169
170     uniques_x[:, 0] = np.nan
171     uniques_y[:, 0] = np.nan
172
173     uniques_x[:, 0] = uniquew
174     uniques_y[:, 0] = uniquew
175     uniques_x[:, 1] = counts
176     uniques_y[:, 1] = counts
177
178     # Store the corresponding values from x and y, vectorized
179     # uniques_x[:, 2:] = np.repeat(x, counts)
180     # uniques_y[:, 2:] = np.repeat(y, counts)
181
182     for i in range(uniquew.shape[0]):
183         num_repeat = int(uniques_x[i, 1])
184         uniques_x[i, 3 : 3 + num_repeat] = x[inverse == i]
185         uniques_y[i, 3 : 3 + num_repeat] = y[inverse == i]
186         mean_of_x = np.mean(uniques_x[i, 3 : 3 + num_repeat], axis=0)
187         mean_of_y = np.mean(uniques_y[i, 3 : 3 + num_repeat], axis=0)
188         uniques_x[i, 2] = mean_of_x
189         uniques_y[i, 2] = mean_of_y
190
191     return uniques_x, uniques_y
192
193
194 def compute_likelihood(theta, data, noise_var):
195     """
196     Compute the likelihood of the data given the parameters theta
197     """
198     noise_mat = np.eye(data.shape[0]) * noise_var
199     evaluate_gaussian(theta, data, noise_var)
200
201
202 def likelihood_mean(theta, time):
203     """
204     Compute the mean of the likelihood function
205     """
206     return theta[0] + theta[1] * time
207
208
209 def likelihood_cov(theta, time):
210     """
211     Compute the covariance of the likelihood function

```

```

212
213     theta: (d, d)
214     time (N, )
215     """
216     # Return array of size (N, d, d)
217
218
219 def compute_confidence_interval(mu, cov):
220     """
221     Compute the confidence interval of the parameters theta
222     """
223     # Compute the likelihood
224     # Compute the posterior
225     # Compute the confidence interval
226     # return theta_low, theta_high
227
228     # sample from a multivariate gaussian 10E5 times
229     samples = mu + np.sqrt(cov) * np.random.randn(int(10e5))
230     evaluations = np.where(np.abs(samples - mu) < 1e-6, 1, 0)
231
232     z = 1.96
233     std_err = np.std(evaluations) / np.sqrt(10e5)
234
235     lower_bound = mu - z * std_err
236     upper_bound = mu + z * std_err
237     return lower_bound, upper_bound
238
239
240 def main():
241     intrinsic_matrix = np.genfromtxt(
242         f"{sys.path[0]}/intrinsic-params.txt", delimiter=" "
243     )
244
245     files = os.listdir(f"{sys.path[0]}/../data/slider_far/split_data")
246     sorted_files = sorted(files)
247
248     # Define data parameters
249     num_data = 7500
250
251     # Load the true parameters that are given from the data in homogenous coordinates
252     true = np.array([0.14, 0, 1])
253     true_pixels = intrinsic_matrix @ true
254
255     # Define prior parameters
256     # X-direction
257     prior_mean_x = np.array([0, 0]).reshape(2, 1)
258     prior_cov_x = np.array([[1, 0], [0, 1]])
259
260     # Y-direction
261     prior_mean_y = np.array([0, 0]).reshape(2, 1)
262     prior_cov_y = np.array([[1, 0], [0, 1]])
263
264     # Array to store the results
265     x_means = np.zeros(shape=(len(sorted_files), 2))
266     x_covs = np.zeros(shape=(len(sorted_files), 2, 2))
267     x_space = np.linspace(0, x_means.shape[0] - 1, x_means.shape[0])
268
269     # fig, axs = plt.subplots(2, 1, figsize=(10, 10))
270     true_x = true_pixels[0]
271     true_y = true_pixels[1]
272
273     y_means = np.zeros(shape=(len(sorted_files), 2))
274     y_covs = np.zeros(shape=(len(sorted_files), 2, 2))
275
276     # Array to store the velocity
277     velocities = np.zeros(shape=(len(sorted_files), 2))

```

```

278
279 num_data = 0
280
281 for file_idx, filename in tqdm(enumerate(sorted_files)):
282
283     # Predicting in the x-direction
284     # Load the data
285     x_values, y_values = store_events_in_matrix(filename)
286
287     training_time = x_values[:, 0].reshape(x_values.shape[0], 1)
288     training_px = x_values[:, 2].reshape(x_values.shape[0], 1)
289
290     xpredictions = np.linspace(-10, 10, 1000)[: , np.newaxis]
291
292     noise_var = np.eye(training_px.shape[0]) * 0.01
293
294     (
295         post_mean_x,
296         post_cov_x,
297         post_pred_pred_mean_x,
298         post_pred_pred_cov_x,
299     ) = linear_regression(
300         training_time,
301         training_px,
302         xpredictions,
303         prior_mean_x,
304         prior_cov_x,
305         noise_var,
306     )
307
308     # Get the mean and covariance of the prediction
309     mean_x = likelihood_mean(post_mean_x, training_time)
310
311     # Predicting in the y-direction
312     training_py = y_values[:, 2].reshape(y_values.shape[0], 1)
313
314     ypredictions = np.linspace(-10, 10, 1000)[: , np.newaxis]
315
316     (
317         post_mean_y,
318         post_cov_y,
319         post_pred_pred_mean_y,
320         post_pred_pred_cov_y,
321     ) = linear_regression(
322         training_time,
323         training_py,
324         ypredictions,
325         prior_mean_y,
326         prior_cov_y,
327         noise_var,
328     )
329
330     # Get the mean and covariance of the prediction
331     mean_y = likelihood_mean(post_mean_y, training_time)
332     cov_y = likelihood_cov(post_cov_y, training_time)
333
334     # Store the results
335     x_means[file_idx] = post_mean_x.reshape(2,)
336     x_covs[file_idx] = post_cov_x
337
338     y_means[file_idx] = post_mean_y.reshape(2,)
339     y_covs[file_idx] = post_cov_y
340
341     # Compute the velocity
342     pixels_x = mean_x[-1] - mean_x[0]
343     pixels_y = mean_y[-1] - mean_y[0]

```



```

344 pixels = np.array([pixels_x[0], pixels_y[0], 1])
345 camera_coordiantes = np.linalg.solve(intrinsic_matrix, pixels)
346 homogeneous_coordinates = camera_coordiantes / camera_coordiantes[-1]
347 x_vel = homogeneous_coordinates[0]
348 y_vel = homogeneous_coordinates[1]
349
350 velocities[file_idx, :] = np.array([x_vel, y_vel]).reshape(2,)
351
352 # Plot the bivariate distribution of the prior
353
354 if file_idx == 0:
355     prior_offset_range_x = np.linspace(
356         prior_mean_x[0][0] - 2.5, prior_mean_x[0][0] + 2.5, 1000
357     )
358     prior_slope_range_x = np.linspace(
359         prior_mean_x[1][0] - 2.5, prior_mean_x[1][0] + 2.5, 1000
360     )
361     prior_offset_range_y = np.linspace(
362         prior_mean_y[0][0] - 2.5, prior_mean_y[0][0] + 2.5, 1000
363     )
364     prior_slope_range_y = np.linspace(
365         prior_mean_y[1][0] - 2.5, prior_mean_y[1][0] + 2.5, 1000
366     )
367
368     offset_range_x = np.linspace(
369         post_mean_x[0][0] - 1000 * post_cov_x[0, 0],
370         post_mean_x[0][0] + 1000 * post_cov_x[0, 0],
371         1000,
372     )
373     slope_range_x = np.linspace(
374         post_mean_x[1][0] - 100 * post_cov_x[1, 1],
375         post_mean_x[1][0] + 100 * post_cov_x[1, 1],
376         100,
377     )
378     offset_range_y = np.linspace(
379         post_mean_y[0][0] - 1000 * post_cov_y[0, 0],
380         post_mean_y[0][0] + 1000 * post_cov_y[0, 0],
381         1000,
382     )
383     slope_range_y = np.linspace(
384         post_mean_y[1][0] - 100 * post_cov_y[1, 1],
385         post_mean_y[1][0] + 100 * post_cov_y[1, 1],
386         100,
387     )
388
389 else:
390
391     prior_offset_range_x = np.linspace(
392         prior_mean_x[0][0] - 3000 * prior_cov_x[0, 0],
393         prior_mean_x[0][0] + 3000 * prior_cov_x[0, 0],
394         3000,
395     )
396     prior_slope_range_x = np.linspace(
397         prior_mean_x[1][0] - 3000 * prior_cov_x[1, 1],
398         prior_mean_x[1][0] + 3000 * prior_cov_x[1, 1],
399         3000,
400     )
401     prior_offset_range_y = np.linspace(
402         prior_mean_y[0][0] - 3000 * prior_cov_y[0, 0],
403         prior_mean_y[0][0] + 3000 * prior_cov_y[0, 0],
404         3000,
405     )
406     prior_slope_range_y = np.linspace(
407         prior_mean_y[1][0] - 3000 * prior_cov_y[1, 1],
408         prior_mean_y[1][0] + 3000 * prior_cov_y[1, 1],
409         3000,

```

```

410     )
411
412     # Plot the bivariate distribution of the posterior
413     offset_range_x = np.linspace(
414         post_mean_x[0][0] - 3000 * post_cov_x[0, 0],
415         post_mean_x[0][0] + 3000 * post_cov_x[0, 0],
416         3000,
417     )
418     slope_range_x = np.linspace(
419         post_mean_x[1][0] - 3000 * post_cov_x[1, 1],
420         post_mean_x[1][0] + 3000 * post_cov_x[1, 1],
421         3000,
422     )
423     offset_range_y = np.linspace(
424         post_mean_y[0][0] - 3000 * post_cov_y[0, 0],
425         post_mean_y[0][0] + 3000 * post_cov_y[0, 0],
426         3000,
427     )
428     slope_range_y = np.linspace(
429         post_mean_y[1][0] - 3000 * post_cov_y[1, 1],
430         post_mean_y[1][0] + 3000 * post_cov_y[1, 1],
431         3000,
432     )
433
434     num_data += x_values.shape[0]
435
436     FIG_DIR = f"{sys.path[0]}/../Report/figures/"
437
438     if file_idx in [0, 49, 99]:
439
440         print("Compute Confidence Intervals:")
441
442         lower_bound_x, upper_bound_x = compute_confidence_interval(
443             post_mean_x[1][0], post_cov_x[1, 1]
444         )
445         lower_bound_y, upper_bound_y = compute_confidence_interval(
446             post_mean_y[1][0], post_cov_y[1, 1]
447         )
448
449         print(
450             f"X-direction: {lower_bound_x} <= {post_mean_x[1][0]} <= {upper_bound_x}"
451         )
452         print(
453             f"Y-direction: {lower_bound_y} <= {post_mean_y[1][0]} <= {upper_bound_y}"
454         )
455
456     # Plot the bivariate distribution of the prior
457     figp, axp = plot_bivariate_gauss(
458         prior_offset_range_x,
459         prior_slope_range_x,
460         prior_mean_x.reshape(2,),
461         prior_cov_x,
462         ax_title="Prior",
463     )
464     figp.suptitle(f"Prior for x-direction after {file_idx} data points")
465     plt.savefig(f"{FIG_DIR}/prior_x-{file_idx}.png")
466     plt.close()
467
468     figp, axp = plot_bivariate_gauss(
469         prior_offset_range_y,
470         prior_slope_range_y,
471         prior_mean_y.reshape(2,),
472         prior_cov_y,
473         ax_title="Prior",
474     )
475     figp.suptitle(f"Prior for y-direction after {num_data} data points")

```

```

476 plt.savefig(f"{FIG_DIR}/prior_y-{{file_idx}}.png")
477 plt.close()
478
479 # Plot the bivariate distribution of the posterior
480 figp, axp = plot_bivariate_gauss(
481     offset_range_x,
482     slope_range_x,
483     post_mean_x.reshape(2,),
484     post_cov_x,
485     ax_title="Posterior",
486 )
487 figp.suptitle(f"Posterior for x-direction after {{num.data}} data points")
488 plt.savefig(f"{FIG_DIR}/posterior_x-{{file_idx}}.png")
489 plt.close()
490
491 figp, axp = plot_bivariate_gauss(
492     offset_range_y,
493     slope_range_y,
494     post_mean_y.reshape(2,),
495     post_cov_y,
496     ax_title="Posterior",
497 )
498 figp.suptitle(f"Posterior for y-direction after {{num.data}} data points")
499 plt.savefig(f"{FIG_DIR}/posterior_y-{{file_idx}}.png")
500 plt.close()
501
502 # Generate 3D plot of the data
503 fig = plt.figure(figsize=(10, 10))
504 fig.suptitle(f"Prediction after {{num.data}} data points")
505
506 ax = fig.add_subplot(211)
507 ax.scatter(training_time, training_px, s=1, label="Events")
508 ax.plot(training_time, mean_x, c="k", label="Predicted trajectory")
509 ax.set_title(r"Regression for x-direction")
510 ax.set_xlabel("Time")
511 ax.set_ylabel("X-direction (pixels)")
512 ax.legend()
513 ax = fig.add_subplot(212)
514 ax.scatter(training_time, training_py, s=1, label="events")
515 ax.plot(training_time, mean_y, c="k", label="Predicted trajectory")
516 ax.set_title(r"Regression for y-direction")
517 ax.set_xlabel("Time")
518 ax.set_ylabel("Y-direction (pixels)")
519 ax.legend()
520 plt.savefig(f"{FIG_DIR}/axes-prediction-{{file_idx}}.png")
521 plt.close()
522 fig = plt.figure()
523 fig.suptitle(f"Prediction after {{num.data}} data points")
524 ax = fig.add_subplot(111, projection="3d")
525 ax.scatter(
526     training_time,
527     training_px,
528     training_py,
529     c="b",
530     marker="o",
531     s=1,
532     label="Events",
533 )
534 ax.set_xlabel("Time")
535 ax.set_ylabel("X-direction")
536 ax.set_zlabel("Y-direction")
537 x_dir = (post_mean_x[0][0] + post_mean_x[1][0] * training_time).reshape(
538     training_time.shape[0],
539 )
540 y_dir = (post_mean_y[0][0] + post_mean_y[1][0] * training_time).reshape(
541     training_time.shape[0],

```

```

542     )
543
544     ax.plot(
545         training_time,
546         x_dir,
547         y_dir,
548         c="r",
549         linewidth=7.0,
550         label="Predicted trajectory",
551     )
552     ax.legend()
553     plt.savefig(f"{FIG_DIR}/3d-prediction_{file_idx}.png")
554     # plt.show()
555     plt.close()
556
557     evaluate_gaussian = lambda x, mean, cov: np.exp(
558         -0.5 * (x - mean) ** 2 / cov
559     )
560
561     fig, axs = plt.subplots(2, 1, figsize=(10, 10))
562     true_x = true_pixels[0]
563     true_y = true_pixels[1]
564     if abs(post_mean_x[1][0] - true_x) > 100:
565         slope_range_x = np.linspace(-200, 200, 1000)
566     else:
567         slope_range_x = np.linspace(
568             post_mean_x[1][0] - 100, post_mean_x[1][0] + 100, 1000
569         )
570     if abs(post_mean_y[1][0] - true_y) > 100:
571         slope_range_y = np.linspace(-200, 200, 1000)
572     else:
573         slope_range_y = np.linspace(
574             post_mean_y[1][0] - 100, post_mean_y[1][0] + 100, 1000
575         )
576     # slope_range_y = np.linspace(-100, 100, 1000)
577     axs[0].plot(
578         slope_range_x,
579         evaluate_gaussian(slope_range_x, post_mean_x[1][0], post_cov_x[1, 1]),
580         label="Prediction on velocity",
581     )
582     axs[1].plot(
583         slope_range_y,
584         evaluate_gaussian(slope_range_y, post_mean_y[1][0], post_cov_y[1, 1]),
585         label="Prediction on velocity",
586     )
587     # Draw a vertical line at the true x
588     axs[0].axvline(true_x, c="k", linestyle="--", label="True velocity")
589     axs[0].axvline(
590         post_mean_x[1][0], c="r", linewidth=1.0, label="Posterior mean"
591     )
592     axs[1].axvline(true_y, c="k", linestyle="--", label="True velocity")
593     axs[1].axvline(
594         post_mean_y[1][0], c="r", linewidth=1.0, label="Posterior mean"
595     )
596     axs[0].set_title(
597         f"Posterior on the velocity for x-direction after {num_data} data points vs  
true parameter"
598     )
599     axs[1].set_title(
600         f"Posterior on the velocity for y-direction after {num_data} data points vs  
true parameter"
601     )
602     axs[1].set_xlabel("Velocity (px/s)")
603     axs[0].set_ylabel("Probability")
604     axs[1].set_ylabel("Probability")
605     axs[0].legend()

```

```

606         axs[1].legend()
607         plt.savefig(f"{FIG_DIR}/posterior_slope-{{file_idx}}-vs-true.png")
608         plt.close()
609     prior_mean_x = post_mean_x
610     prior_cov_x = post_cov_x
611
612     prior_mean_y = post_mean_y
613     prior_cov_y = post_cov_y
614
615     fig, axs = plt.subplots(2, 1, figsize=(10, 10))
616     axs[0].axhline(true_x, c="b", linestyle="--", label="True velocity")
617     axs[1].axhline(true_y, c="b", linestyle="--", label="True velocity")
618     axs[0].plot(x_space, x_means[:, 1], c="k", label="Predicted vel in x")
619     axs[1].plot(x_space, y_means[:, 1], c="k", label="Predicted vel in y")
620     axs[0].set_ylabel("Velocity (px/s)")
621     axs[1].set_xlabel("Data Segments")
622     axs[1].set_ylabel("Velocity (px/s)")
623     fig.suptitle("Convergence Predicted Velocity over Data Segments")
624     axs[0].legend()
625     axs[1].legend()
626     plt.savefig(f"{FIG_DIR}/predicted_velocity_over_data_segments.png")
627     plt.show()
628
629
630 if __name__ == "__main__":
631
632     main()

```

1.2.2 Plotting and Toy Problems

```

1  import matplotlib.pyplot as plt
2  from mpl_toolkits.mplot3d import Axes3D
3
4  class EventPlotter:
5      '''
6      Class for plotting any event that is input
7      '''
8
9      def __init__(self, t, x, y, pol = None):
10         self.x = x
11         self.y = y
12         self.time = t
13         if pol is not None:
14             self.pol = pol
15             self.pol_exists = True
16         else:
17             self.pol_exists = False
18
19
20     def plot_events(self, t_start=0, t_end=-1):
21         '''
22         Plot the events
23         '''
24         fig = plt.figure()
25         ax = fig.add_subplot(111, projection='3d')
26         ax.scatter(self.time[t_start:t_end], self.x[t_start:t_end], self.y[t_start:t_end])
27         ax.set_xlabel('Time')
28         ax.set_ylabel('X')
29         ax.set_zlabel('Y')
30         plt.show()
31
32
33     def plot_x_axis(self, t_start=0, t_end=-1):
34         '''
35         Plot the events only on the x-axis
36         '''
37         fig = plt.figure()

```

```

38         ax = fig.add_subplot(111)
39         ax.scatter(self.time[t_start:t_end], self.x[t_start:t_end], s=1)
40         ax.set_xlabel('X')
41         ax.set_ylabel('Time')
42         plt.show()
43
44     def plot_y_axis(self, t_start=0, t_end=-1):
45         '''
46         Plot the events only on the y-axis
47         '''
48         fig = plt.figure()
49         ax = fig.add_subplot(111)
50         ax.scatter(self.time[t_start:t_end], self.y[t_start:t_end], s=1)
51         ax.set_xlabel('Y')
52         ax.set_ylabel('Time')
53         plt.show()

1  import numpy as np
2  import matplotlib.pyplot as plt
3
4  def normpdf(x, mean, cov):
5      n, N = x.shape
6      preexp = 1.0 / (2.0 * np.pi)**(n/2) / np.linalg.det(cov)**0.5
7      diff = x - np.tile(mean[:, np.newaxis], (1, N))
8      sol = np.linalg.solve(cov, diff)
9      inexp = np.einsum("ij,ij->j", diff, sol)
10     out = preexp * np.exp(-0.5 * inexp)
11     return out
12
13 def eval_normpdf_on_grid(x, y, mean, cov):
14     XX, YY = np.meshgrid(x,y)
15     pts = np.stack((XX.reshape(-1), YY.reshape(-1)), axis=0)
16     evals = normpdf(pts, mean, cov).reshape(XX.shape)
17     return XX, YY, evals
18
19 def plot_bivariate_gauss(x, y, mean, cov, ax_title=None):
20     std1 = cov[0,0]**0.5
21     std2 = cov[1,1]**0.5
22     mean1 = mean[0]
23     mean2 = mean[1]
24     XX, YY, evals = eval_normpdf_on_grid(x, y, mean, cov)
25     fig, axis = plt.subplots(2,2, figsize=(10,10))
26     axis[0,0].plot(x, normpdf(x[np.newaxis,:], np.array([mean1]), np.array([[std1**2]])))
27     axis[0,0].set_ylabel(r'$f_{X_1}$')
28     axis[0,0].set_title(f'{ax_title} distribution on \nthe initial position parameter')
29     axis[1,1].plot(normpdf(y[np.newaxis,:], np.array([mean2]), np.array([[std2**2]])), y)
30     axis[1,1].set_xlabel(r'$f_{X_2}$')
31     axis[1,1].set_title(f'{ax_title} distribution on \nthe velocity parameter')
32     axis[1,0].contourf(XX, YY, evals)
33     axis[1,0].set_xlabel(r'$x_1$')
34     axis[1,0].set_ylabel(r'$x_2$')
35     axis[0,1].set_visible(False)
36     return fig, axis

1  import numpy as np
2  import matplotlib.pyplot as plt
3  from mpl_toolkits.mplot3d import Axes3D
4  import math
5
6
7  class ToyEvent:
8      """
9      Class for generating toy problems.
10     """
11
12     def __init__(self, timespan, dt, x_parameter, y_parameter, noise, shape='dot', num_pts
        =1, offtimed=False):

```

```

13         '''
14         Initialize the toy problem class
15         '''
16         self.timespan = timespan
17         self.dt_ = dt
18         self.dt = dt if type(dt) == int else int(dt)
19         self.t = int(math.ceil(timespan/dt))
20
21         self.tmp_x = []
22         self.tmp_y = []
23
24         self.data_dim = 10
25
26         self.x_parameter = np.ones(shape=(num_pts*self.data_dim)) * x_parameter
27         self.y_parameter = y_parameter
28         self.shape = shape
29         self.offtimed = offtimed
30         self.noise = noise
31         self.num_pts = num_pts
32         tmp_time = np.arange(0, timespan, dt).reshape(self.t, 1)
33         self.time = tmp_time
34
35
36         # Stack if there are multiple objects to track
37         for i in range((num_pts*self.data_dim)-1):
38             self.time = np.hstack((self.time, tmp_time))
39         self.initial_points_x = np.ones(shape=(
40             self.t, self.num_pts*self.data_dim)) * 0 #@ (5*np.random.randn(self.num_pts*self
41             .data_dim, self.num_pts*self.data_dim))
42         self.initial_points_y = np.ones(shape=(
43             self.t, self.num_pts * self.data_dim)) * 0 #@ (5*np.random.randn(self.num_pts*
44             self.data_dim, self.num_pts*self.data_dim))
45         self.x = np.zeros(shape=(self.t, self.num_pts * self.data_dim))
46         self.y = np.zeros(shape=(self.t, self.num_pts * self.data_dim))
47
48         # print(self.x.shape)
49
50         self.generate_toy_problem()
51
52 def generate_toy_problem(self):
53     """
54     Generates a toy problem.
55     :return:
56     """
57     if self.shape == 'dot' and self.offtimed == False:
58         return self.generate_dot_problem()
59     elif self.shape == 'dot' and self.offtimed:
60         return self.generate_offtimed_dot_problem()
61     elif self.shape == 'circle':
62         return self.generate_circle_problem()
63     else:
64         raise ValueError('Shape not recognized.')
65
66 def generate_dot_problem(self):
67     """
68     Generate a multiple dot with noise
69     """
70     x = self.x_parameter * self.time
71     x_noise = self.noise * np.random.randn(self.t, self.num_pts * self.data_dim)
72     self.x = x + x_noise + self.initial_points_x
73
74     y = self.y_parameter * self.time
75     y_noise = self.noise * np.random.randn(self.t, self.num_pts * self.data_dim)
76     self.y = y + y_noise + self.initial_points_y
77
78 def get_number_of_repeating_points(self, arr):

```

```

77     '''
78     Returns the number of repeating points in an np.array
79     '''
80     return len(np.unique(arr, axis=0))
81
82 def get_nearest_value(self, arr1, arr2):
83     '''
84     Returns the nearest value in arr2 to the value in arr1
85     '''
86     # return (np.abs(arr1 - arr2)).argmin(axis=1)
87     return np.array([np.argmin(np.abs(arr2 - x)) for x in arr1])
88
89 def generate_offtimed_dot_problem(self):
90     '''
91     Generate a multiple dot with noise
92     '''
93     time = 0
94     events = np.array([[0, 0, time]])
95     self.tmp_x = self.x
96     self.tmp_y = self.y
97     while time < self.timespan:
98         num_events = int(5 + 100*np.random.randn())
99         num_events = -1*num_events if num_events < 0 else num_events
100         for i in range(num_events):
101             x = np.mean(self.x_parameter) * time + self.noise * np.random.randn()
102             y = np.mean(self.y_parameter) * time + self.noise * np.random.randn()
103             events = np.vstack((events, [x, y, time]))
104             # add x to tmp_x with the corresponding time index
105             time_idx = np.where(self.time == time)
106             # print(time)
107             # self.tmp_x[np.where(self.time == time)] += x
108             # self.tmp_y[np.where(self.time == time)] += y
109             time += self.dt_
110
111     self.x = events[:, 0].reshape(events[:, 0].shape[0], 1)
112     self.y = events[:, 1].reshape(events[:, 1].shape[0], 1)
113     self.time = events[:, 2].reshape(events[:, 2].shape[0], 1)
114
115
116 def ravel_events(self):
117     '''
118     Ravel the events into a single array
119     '''
120     if self.offtimed == False: raise ValueError('Not implemented')
121     # find the indices of the repeating values in the time array

```